

CS 5200

Database Systems

The Relational Model

References:

Database Management Systems, by Ramakrishnan and Gehrke

(Chapter 3, Sections 3.1-3.5 of the Textbook)



Hazra Imran

Structure of Relational Databases

- *Relational database*: a set of *relations*
- *Relation*: made up of 2 parts:
 - *Schema* : specifies name of relation, plus name and domain (type) of each attribute.
 - e.g., Student (*sid*: string, *name*: string, *address*: string, *phone*: string, *major*: string).
 - *Instance* : a *table*, with rows and columns.
#Rows = cardinality
#Columns = arity / degree
- *Relational Database Schema*: collection of schemas in the database
- *Database Instance*: a collection of instances of its relations

Example of a Relation Instance

Student

sid	name	address	phone	major
5345	M. Jones	6789 W. 12 th Ave., Vancouver	333-4444	CPSC
2345	M. Smith	2020 Drake Street, Burnaby	109-2222	MATH
1400	B. Smith	3789 No 3 Richmond	222-1234	CPSC
1653	S. Fang	null	null	null

Example of a Relation Instance

relation name → **Student**

attribute, column name

tuple, row, record →

domain value →

sid	name	address	phone	major
5345	M. Jones	6789 W. 12 th Ave., Vancouver	333-4444	CPSC
2345	M. Smith	2020 Drake Street, Burnaby	109-2222	MATH
1400	B. Smith	3789 No 3 Richmond	222-1234	CPSC
1653	S. Fang	null	null	null

- degree/arity = 5; Cardinality = 4,
- Order of rows isn't important
- Order of attributes isn't important (except in some query languages)

Formal Structure

- Formally, a relation r is a set (a_1, a_2, \dots, a_n) where a_i is in D_i , the domain (set of allowed values) of the i -th attribute.
- Attribute values are atomic, i.e., integers, floats, strings
- A domain contains a special value *null* indicating that the value is not known.
- If A_1, \dots, A_n are attributes with domains D_1, \dots, D_n , then $(A_1:D_1, \dots, A_n:D_n)$ is a *relation schema* that defines a relation type - sometimes we leave off the domains

Example of a formal definition

Student

sid	name	address	phone	major
5345	M. Jones	6789 W. 12 th Ave., Vancouver	333-4444	CPSC
2345	M. Smith	2020 Drake Street, Burnaby	109-2222	MATH
1400	B. Smith	3789 No 3 Richmond	222-1234	CPSC
1653	S. Fang	null	null	null

Relational Schema

Student(sid: integer, name: string, address: string, phone: string, major: string)

Or, without the domains:

Student (sid, name, address, phone, major)

Relational Query Languages

- A major strength of the relational model: simple, powerful *querying* of data.
- Queries can be written intuitively; DBMS is responsible for efficient evaluation.
 - Precise semantics for relational queries.
 - Allows optimizer to extensively re-order operations, while ensuring that the answer does not change.

The SQL Query Language

- Developed by IBM (System R) in the 1970s
- Standards:
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision, current standard)
 - SQL-99 (major extensions)
 - ...SQL-2003, 2008, 2011, 2016

<http://www.jcc.com/resources/sql-standards>)



Raymond Boyce



Don Chamberlain

A glance at the SQL Query Language (1/2)

- Find the id, names and phones of all CPSC students:

```
SELECT sid, name, phone
FROM Students
WHERE major="CPSC"
```

sid	name	phone
5345	M. Jones	333-4444
1400	B. Smith	222-1234

Student

sid	name	address	phone	major
5345	M. Jones	6789 W. 12 th Ave., Vancouver	333-4444	CPSC
2345	M. Smith	2020 Drake Street, Burnaby	109-2222	MATH
1400	B. Smith	3789 No 3 Richmond	222-1234	CPSC
1653	S. Fang	null	null	null

- To select whole rows , replace “SELECT sid, name, phone ” with “SELECT * ”

A glance at the SQL Query Language (2/2): Querying Multiple Tables

Student

sid	name	address	phone	major
5345	M. Jones	CPSC
...

Grade

sid	dept	Course#	marks
5345	CPSC	354	86
...

- To select name and marks of the students who have taken some CPSC course, we write:

SELECT name, marks

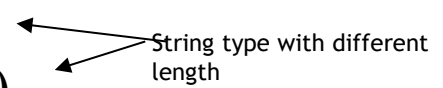
FROM Student, Grade

WHERE Student.sid = Grade.sid AND dept = 'CPSC'

Creating Relations in SQL/DDL

- Creates the Student relation
 - the type (**domain**) of each attribute is specified
 - enforced when tuples are added or modified
- The Grade table holds information about courses that a students takes

```
CREATE TABLE Student
(sid INTEGER,
 name CHAR(20),
 address CHAR(30),
 phone CHAR(13) DEFAULT '999999999',
 major CHAR(4))
```



sid	name	address	phone	major
5345	M. Jones	CPSC
...

```
CREATE TABLE Grade
(sid INTEGER,
 dept CHAR(4),
 course# CHAR(3),
 mark INTEGER)
```

sid	dept	Course#	mark
5345	CPSC	354	86
...

Deleting Table

To delete a table use the DROP TABLE statement

DROP TABLE Student

- Destroys the relation Student.
- Schema information *and* tuples are deleted.

Altering Table

- Columns can be added or removed to tables using the **ALTER TABLE** statement
 - ADD to add a column and
 - DROP to remove a column
- The schema of Students is altered by adding a new attribute; every tuple in current instance is extended with a *null* value in the new attribute.

```
ALTER TABLE Student  
ADD gpa REAL;
```

```
ALTER TABLE Student  
DROP COLUMN gpa;
```

Null Values

- A field can have a special value NULL.
- Null can many things.
 - Field is unknown or missing or unavailable
 - Field may not apply in certain cases.
- Not allowing NULL in some cases is a way to enforce total participation.

Adding Tuples

- Can insert a single tuple using:

```
INSERT INTO Student (sid, name, address, phone, major)
VALUES ('5345', 'M. Currie', '1235 Burrad St', '882-4444', 'PHYS')
```

sid	name	address	phone	major
5345	M. Currie	12345 Burrad St	882-4444	PHYS
1234	Smith	234 Kingsway	567-9878	CS
3453	Smith	345 Terminal Ave	567-4532	MATH
...

The list of column names is optional

- If omitted the values must be in the same order as the columns

Deleting Tuples

Can delete all tuples satisfying some condition (e.g., name = 'Smith'):

```
DELETE
FROM   Student
WHERE  name = 'Smith'
```

sid	name	address	phone	major
5345	M. Currie	12345 Burrad St	882-4444	PHYS
1234	Smith	234 Kingsway	567-9878	CS
3453	Smith	345 Terminal Ave	567-4532	MATH
...

→ *Powerful variants of these commands exist; more later*

Update Tuples

- Use the UPDATE statement to modify a record, or records, in a table
- Note that the WHERE statement is evaluated before the SET statement
- Like DELETE the WHERE clause specifies which records are to be updated

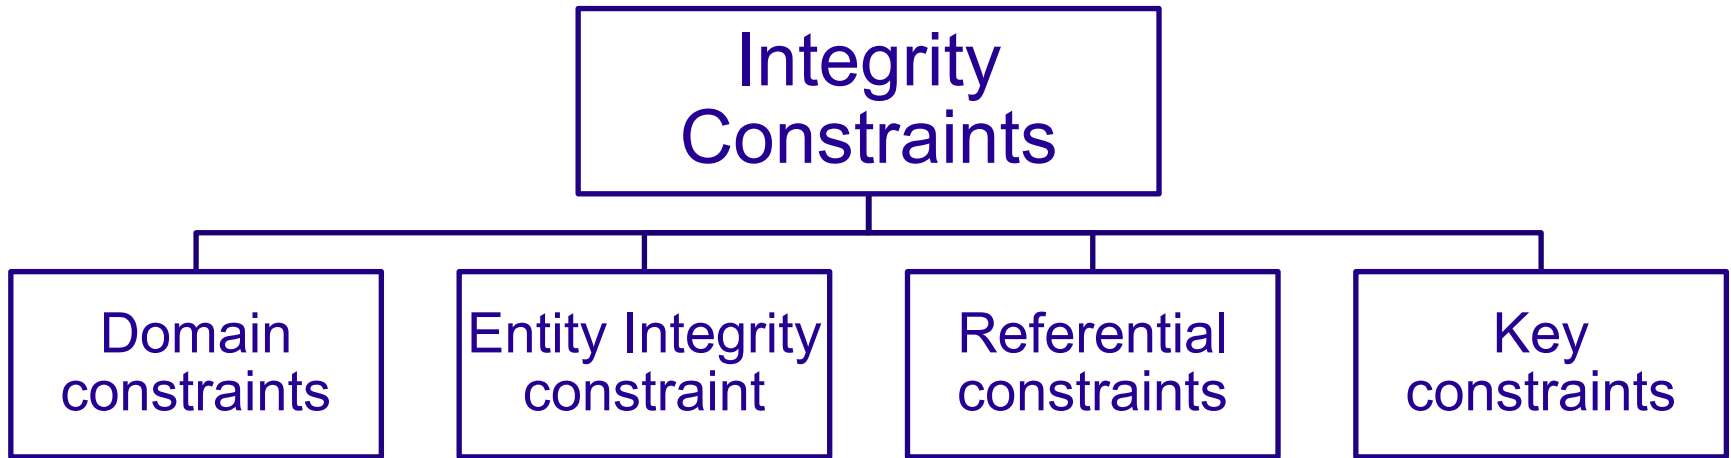
```
UPDATE STUDENT  
SET address = '333 Argyle Street'  
WHERE name = 'M. Currie'
```

Integrity Constraints (ICs)

“Integrity is doing the right thing, even when no one is watching” (By : C.S. Lewis)

- **IC**: condition that must be true for *any* instance of the database; e.g., domain constraints
 - ICs are specified when schema is defined
 - ICs are checked when relations are modified
- A *legal* instance of a relation is one that satisfies all specified ICs
 - DBMS should not allow illegal instances
 - Avoids data entry errors, too!
- The types of IC's depend on the data model.
 - Next up: constraints for relational databases

IC



Candidate Key

- A set of fields is a **superkey** for a relation if :
 - No two distinct tuples can have same values in all key fields.
- **Candidate key** = minimal superkey = no subset of fields is a superkey.
- One of the possible keys is chosen (by the DBA) to be the **primary key** (PK).

```
CREATE TABLE Student
(sid          INTEGER,
 name        CHAR(20),
 address     CHAR(30),
 phone       CHAR(13),
 major       CHAR(4))
```

- e.g.
 - {*sid*, *name*} is a superkey
 - *sid* is the primary key for Students

PRIMARY KEY in SQL

- A **PRIMARY KEY** constraint specifies a table's primary key
 - values for primary key must be unique
 - a primary key attribute cannot be *null*
- ```
CREATE TABLE Student
(sid INTEGER PRIMARY KEY,
name CHAR(20),
address CHAR(30),
phone CHAR(13),
major CHAR(4))
```
- Other keys are specified using the **UNIQUE** constraint.
    - The values for a group of attributes must be unique (if they are not null).
    - These attributes can be null
  - **Key constraints** are checked when
    - new values are inserted
    - values are modified

# Super key, Candidate key and Primary key

---

## Students

| <u>sid</u> | name | address | phone | major |
|------------|------|---------|-------|-------|
|            |      |         |       |       |
|            |      |         |       |       |
|            |      |         |       |       |
|            |      |         |       |       |

# PRIMARY KEY in SQL (cont')

(Ex.1- Normal) “For a given student and course, there is a single mark’s value.” **vs**

```
CREATE TABLE Grade
(sid INTEGER,
 dept CHAR(4),
 course# CHAR(3),
 mark INTEGER,
 PRIMARY KEY (sid,dept,course#))
```

(Ex.2 - Bad) “Students can take a course once, and receive a single grade for that course; further, no two students in a course receive the same grade.”

```
CREATE TABLE Grade2
(sid INTEGER,
 dept CHAR(4),
 course# CHAR(3),
 marks INTEGER,
 PRIMARY KEY(sid,dept,course#),
 UNIQUE (dept,course#,mark))
```

## Grade

| <u>sid</u> | <u>dept</u> | <u>Course#</u> | mark |
|------------|-------------|----------------|------|
|            |             |                |      |
|            |             |                |      |
|            |             |                |      |

For single attribute keys, can also be declared on the same line as the attribute

# Foreign Keys Constraints

- **Foreign key** : Set of attributes in one relation used to 'reference' a tuple in another relation.
  - Must correspond to the primary key of the other relation.
  - Like a 'logical pointer'.
- E.g.: Grade(*sid*, dept, course#, marks)
  - *sid* is a foreign key referring to **Student**:
  - (*dept*, *course#*) is a foreign key referring to **Course**

| Student    |      |         |       |       |
|------------|------|---------|-------|-------|
| <u>SID</u> | Name | Address | Phone | Major |

| Grade      |             |                |       |
|------------|-------------|----------------|-------|
| <u>SID</u> | <u>Dept</u> | <u>course#</u> | Marks |

| Course      |                |       |         |
|-------------|----------------|-------|---------|
| <u>Dept</u> | <u>Course#</u> | cname | credits |

## Referential Integrity

- All foreign keys reference existing entities.
  - i.e. there are no dangling references
  - all foreign key constraints are enforced



# Foreign Keys in SQL

- Only students listed in the Student relation should be allowed to have grades for courses that are listed in the Course relation.

CREATE TABLE Grade

(sid INTEGER, dept CHAR(4), course# CHAR(3), marks INTEGER,

PRIMARY KEY (sid,dept,course#),

FOREIGN KEY (sid) REFERENCES Student(sid),

FOREIGN KEY (dept, course#) REFERENCES Course(dept, cnum))

Primary key in Student

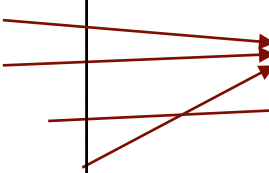
Primary key in Course

Grade

| sid  | dept | course# | marks |
|------|------|---------|-------|
| 5345 | CPSC | 101     | 80    |
| 5345 | RELG | 100     | 45    |
| 1400 | MATH | 200     | null  |
| 5345 | HIST | 201     | 60    |

Student

| sid  | name     | address | Phone | major |
|------|----------|---------|-------|-------|
| 5345 | M. Jones | ....    | ...   | ...   |
| 1400 | B. Smith | ....    | ...   | ...   |
| 2354 | M. Smith | ....    | ...   | ...   |



# Self Referencing Relations

Goal: have managerID be foreign key reference for **same table** Emps.

| id | sin  | name | managerID |
|----|------|------|-----------|
| 1  | 1000 | John | Null      |
| 2  | 1001 | Jack | 1         |

Could foreign key be null?

For **referential integrity** to hold in a relational database, any field in a table that is declared a foreign key should contain either a **null value**, or only values from a parent table's **primary key**.

# Enforcing Referential Integrity

- *sid* in Grade is a foreign key that references Student.
- What should be done if a Grade tuple with a non-existent student id is inserted?
- What should be done if a **Student tuple** is deleted , e.g. *sid* = 5345? *(Reject it!)*
  - Also delete all Grade tuples that refer to it?
  - Disallow deletion of this particular Student tuple?
  - Set *sid* in Grade tuples that refer to it, to *null*, (the special value denoting 'unknown' or "inapplicable".)
    - problem if *sid* is part of the primary key
  - Set *sid* in Grade tuples that refer to it, to a *default sid*.
- Similar if primary key of a Student tuple is updated

Grade

| <u>sid</u> | <u>dept</u> | <u>course#</u> | marks |
|------------|-------------|----------------|-------|
| 5345       | CPSC        | 101            | 80    |
| 5345       | RELG        | 100            | 45    |
| 1400       | MATH        | 200            | null  |
| 5345       | HIST        | 201            | 60    |

Student

| <u>sid</u> | name     | address | Phone | major |
|------------|----------|---------|-------|-------|
| 5345       | M. Jones | ....    | ...   | ...   |
| 1400       | B. Smith | ....    | ...   | ...   |
| 2354       | M. Smith | ....    | ...   | ...   |

# Referential Integrity in SQL/92

---

- SQL/92 supports all 4 options on deletes and updates.
- Default is **NO ACTION**  
(*delete/update is rejected*)
- **CASCADE** (also updates/deletes all tuples that refer to the updated/deleted tuple)
- **SET NULL / SET DEFAULT**  
(referencing tuple value is set to the default foreign key value )

```
CREATE TABLE Grade
(sid CHAR(8),
 dept CHAR(4),
 course# CHAR(3),
 marks INTEGER,
 PRIMARY KEY (sid,dept,course#),
 FOREIGN KEY (sid) REFERENCES Student
 ON DELETE CASCADE
 ON UPDATE CASCADE
 FOREIGN KEY (dept, course#)
 REFERENCES Course
 ON DELETE SET DEFAULT
 ON UPDATE CASCADE);
```

Oracle does not support “on update”

# Where do ICs Come From?

---

- ICs are based upon the real-world semantics being described (in the database relations).
- We *can* check a database instance to verify an IC, but we *cannot* tell the ICs by looking at the instance.
  - For example, even if all student names differ, we cannot assume that name is a key.
  - An IC is a statement about *all possible* instances.
- All constraints must be identified during the conceptual design.
- Some constraints can be explicitly specified in the conceptual model
- Others are written in a more general language.

# Mapping of ERD to relation model

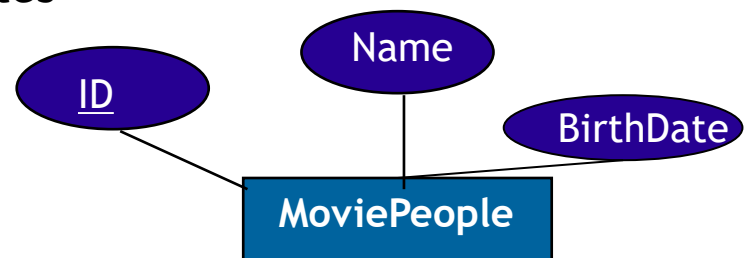
---

# Logical DB Design: ER to Relational

Entity sets to tables.

- Each entity set is mapped to a table.
  - entity attributes become table attributes
  - entity keys become table keys

```
CREATE TABLE MoviePeople
 (ID CHAR(11),
 Name CHAR(20),
 BirthDate Date,
 PRIMARY KEY (ID))
```



MoviePeople

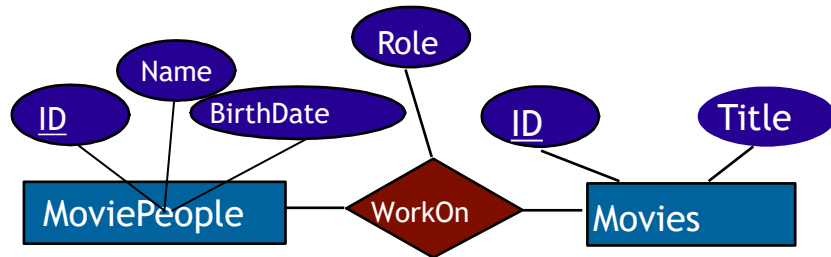
| <u>ID</u> | Name           | BirthDate  |
|-----------|----------------|------------|
| MP001     | Chris Columbus | 10/09/1958 |
| MP002     | Tom Hardy      | 15/09/1977 |
| MP003     | Kate Winslet   | 05/10/1975 |

Movie

| <u>ID</u> | title        |
|-----------|--------------|
| 1         | Harry Potter |
| 2         | The Drop     |
| 4         | Titanic      |

# Relationship Sets to Tables

- A relationship set is mapped to a single relation (table).
- Simple case: relationship has no constraints (i.e. many-to-many)
  - In this case, attributes of the table must include:
    - Keys for each participating entity set as foreign keys. (This is a *key* for the relation)
    - All descriptive attributes.



```
CREATE TABLE WorkOn(
 PID CHAR(11),
 MID INTEGER,
 Role CHAR(20),
 PRIMARY KEY (PID, MID),
 FOREIGN KEY (PID) REFERENCES MoviePeople,
 FOREIGN KEY (MID) REFERENCES Movies)
```

PID and MID CANNOT be null  
and were renamed

WorksOn

| <b>PID</b> | <b>MID</b> | Role           |
|------------|------------|----------------|
| MP002      | 2          | Bob Saginowski |
| MP003      | 4          | Rose DeWitt    |

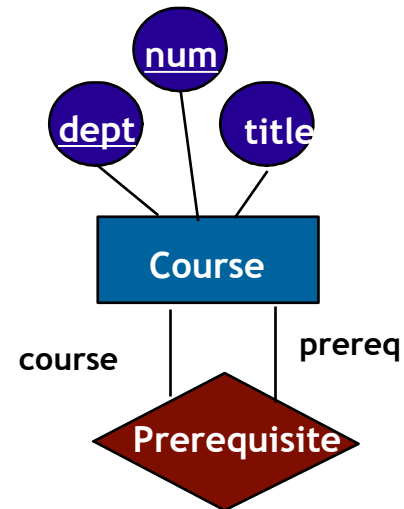
Bold column means FK



# Relationship Sets to Tables (cont')

- In some cases, we need to use the roles:

```
CREATE TABLE Prerequisite(
 course_dept CHAR(4),
 course_num CHAR(3),
 prereq_dept CHAR(4),
 prereq_num CHAR(3),
 PRIMARY KEY (course_dept, course_num, prereq_dept, prereq_num),
 FOREIGN KEY (course_dept, course_num) REFERENCES Course(dept, num),
 FOREIGN KEY (prereq_dept, prereq_num) REFERENCES Course(dept, num))
```



Prerequisite

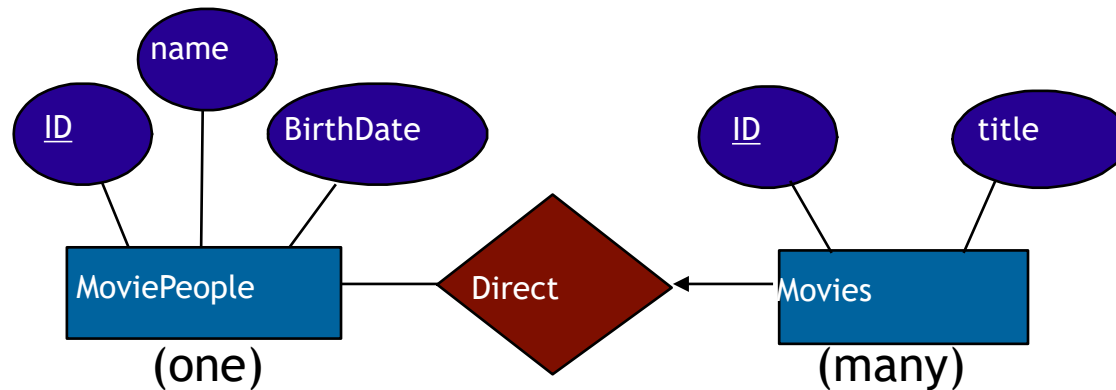
| <u>Course_dept</u> | <u>Course_num</u> | <u>prereq_dept</u> | <u>prereq_num</u> |
|--------------------|-------------------|--------------------|-------------------|
| CPSC               | 304               | CPSC               | 221               |
| CPSC               | 304               | EECE               | 320               |
| ....               | ...               | ....               | ...               |

Course

| <u>dept</u> | <u>num</u> | title                              |
|-------------|------------|------------------------------------|
| CPSC        | 304        | Database Systems                   |
| CPSC        | 221        | Basic Algo and DS                  |
| EECE        | 320        | Discrete Structures and Algorithms |
| ....        | ...        | ....                               |

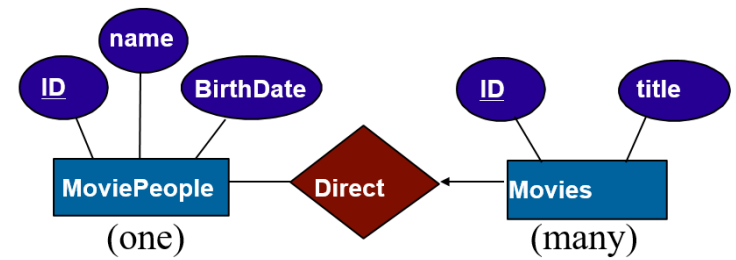
# Review: Key Constraints

- Each movie has at most one director, according to the key constraint on Direct.



- How can we take advantage of this?

# Translating ER Diagrams with Key Constraints



- Method 1 (unsatisfactory):
  - Create a separate table for Direct:
  - Note that **MID** is the key now!
  - Create separate tables for MoviePeople and Movies.
- Method 2 (better)
  - Since each movie has a unique director, we can **combine Direct and Movies into one table**.
  - Create another table for MoviePeople

```

CREATE TABLE Direct(
 MID INTEGER,
 PID CHAR(11),
 PRIMARY KEY (MID),
 FOREIGN KEY (PID) REFERENCES MoviePeople,
 FOREIGN KEY (MID) REFERENCES Movies)

```

Vs.

```

CREATE TABLE Directed_Movie(
 MID INTEGER,
 title CHAR(20),
 PID CHAR(11),
 PRIMARY KEY (MID),
 FOREIGN KEY (PID) REFERENCES MoviePeople
 ON DELETE SET NULL
 ON UPDATE CASCADE)

```

MoviePeople

| <u>ID</u> | Name           |
|-----------|----------------|
| MP001     | Chris Columbus |
| MP002     | Tom Hardy      |
| MP003     | Kate Winslet   |

Movie

| <u>ID</u> | title        |
|-----------|--------------|
| 1         | Harry Potter |
| 2         | The Drop     |
| 4         | Titanic      |

Direct

| <u>MID</u> | PID   |
|------------|-------|
| 1          | MP001 |
| 2          | Null  |
| 4          | Null  |

Directed\_Movie

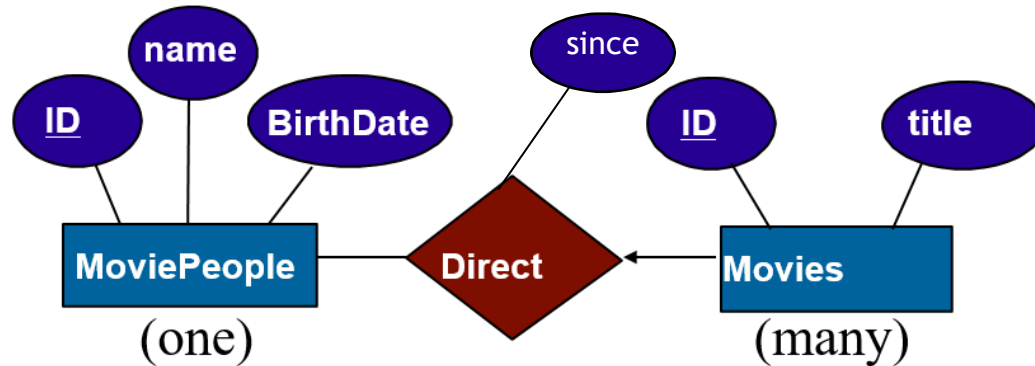
| <u>ID</u> | title        | PID   |
|-----------|--------------|-------|
| 1         | Harry Potter | MP001 |
| 2         | The Drop     | Null  |
| 4         | Titanic      | Null  |

MoviePeople

| <u>ID</u> | Name           |
|-----------|----------------|
| MP001     | Chris Columbus |
| MP002     | Tom Hardy      |
| MP003     | Kate Winslet   |

What if Chris is deleted from MoviePeople?

# If relationship has descriptive attributes then?



**Directed\_Movie**

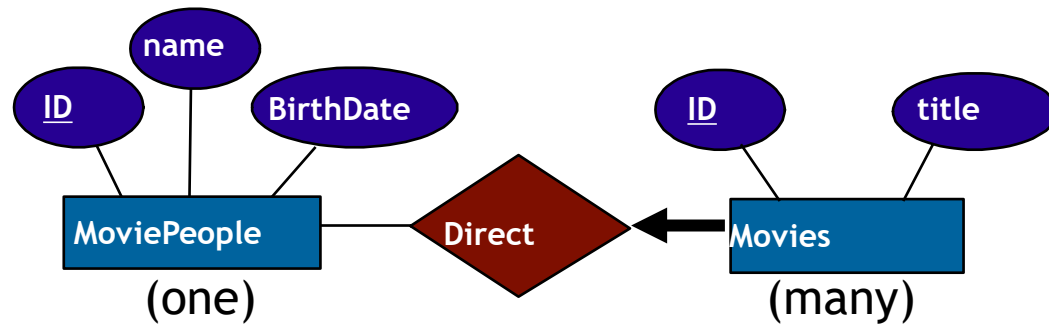
| <u>ID</u> | title        | <b>PID</b> | Since      |
|-----------|--------------|------------|------------|
| 1         | Harry Potter | MP001      | 4 Jan 2018 |
| 2         | The Drop     | NULL       | NULL       |
| 4         | Titanic      | NULL       | NULL       |

**MoviePeople**

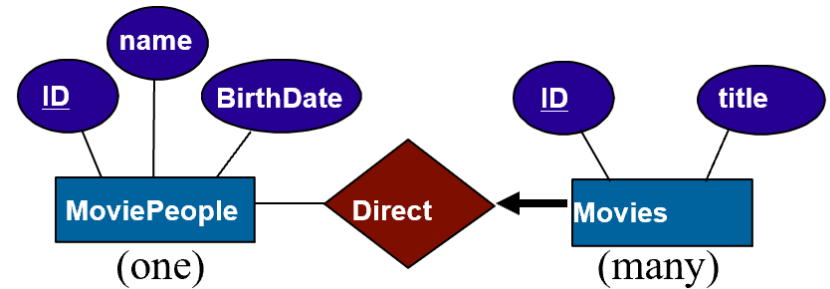
| <u>ID</u> | Name           |
|-----------|----------------|
| MP001     | Chris Columbus |
| MP002     | Tom Hardy      |
| MP003     | Kate Winslet   |

# What if “Every Movie must have a director”

---



# What if “Every Movie must have a director” directed\_movie table



```

CREATE TABLE Direct(
 MID INTEGER,
 PID CHAR(11),
 PRIMARY KEY (MID),
 FOREIGN KEY (PID) REFERENCES MoviePeople,
 FOREIGN KEY (MID) REFERENCES Movies)

```

Vs.

```

CREATE TABLE Directed_Movie(
 MID INTEGER,
 title CHAR(20),
 PID CHAR(11) NOT NULL,
 PRIMARY KEY (MID),
 FOREIGN KEY (PID) REFERENCES MoviePeople
 ON DELETE NO ACTION
 ON UPDATE CASCADE)

```

A MoviePeople tuple cannot be deleted if it is pointed to by a directed\_movie tuple

MoviePeople

| <u>ID</u> | Name           |
|-----------|----------------|
| MP001     | Chris Columbus |
| MP002     | Tom Hardy      |
| MP003     | Kate Winslet   |

Movie

| <u>ID</u> | title        |
|-----------|--------------|
| 1         | Harry Potter |
| 2         | The Drop     |
| 4         | Titanic      |

Direct

| <u>MID</u> | PID   |
|------------|-------|
| 1          | MP001 |
| 2          | Null  |
| 4          | Null  |

Directed\_Movie

| <u>ID</u> | title        | PID   |
|-----------|--------------|-------|
| 1         | Harry Potter | MP001 |

MoviePeople

| <u>ID</u> | Name           |
|-----------|----------------|
| MP001     | Chris Columbus |
| MP002     | Tom Hardy      |
| MP003     | Kate Winslet   |

What if Chris is deleted from MoviePeople?

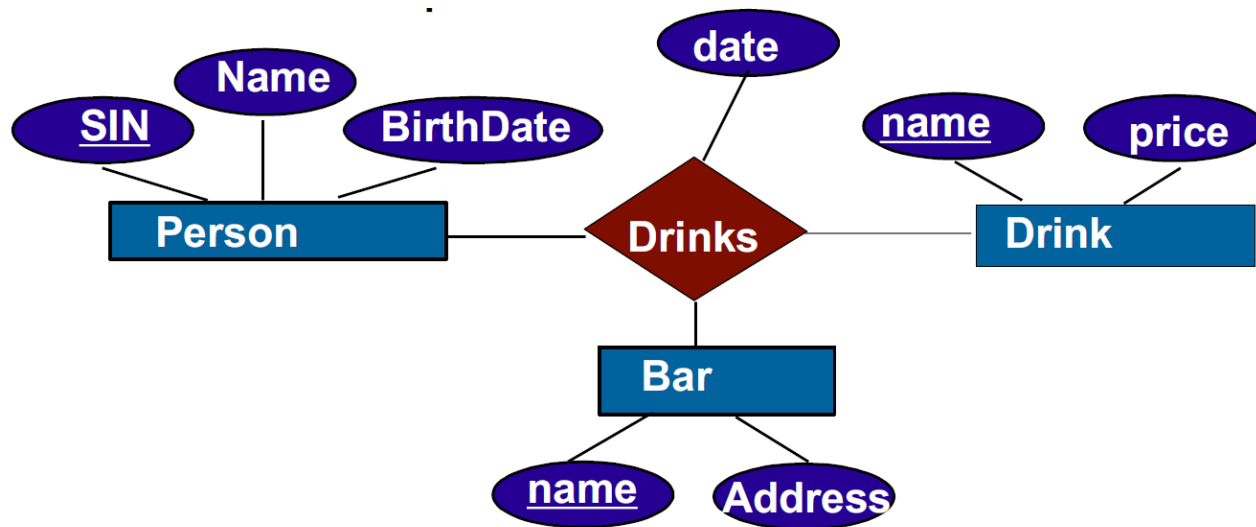
# Summary

---

- When there is no cardinality constraint (i.e., M:M relationship), one table per participating entity set and one separate table for relationship set.
- When there is a cardinality constraint (e.g., many-to-one),
  - one table for the “one” side (e.g., MoviePeople)
  - one table for the “many” side which also captures the relationship (e.g. Movies + Direct)

# Key constraints on non-binary relationships

- Ternary relation Drinks relates entity sets Person, Bar and Drink, and has descriptive attribute *date*.





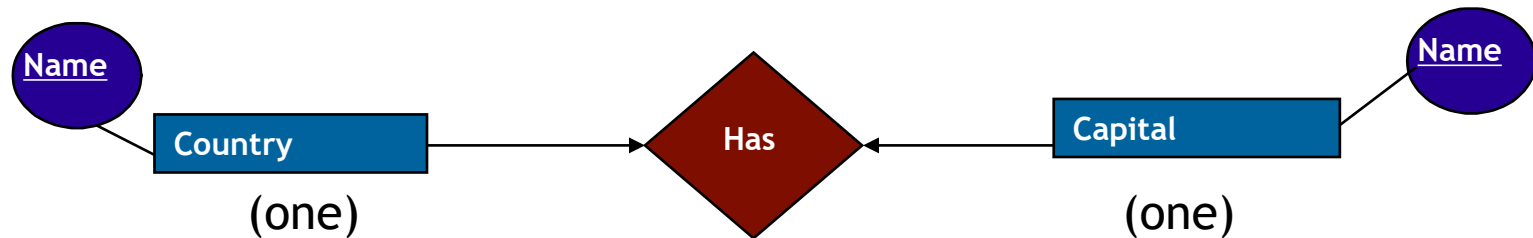
# We'd covered basic translating of ER to relational

---

- Short version: everything's a table
- Slightly longer version:
  - In many to many relationships, create one table per entity and one table per relationship. Link the two by foreign keys
  - In 1:M, merge the relationship with the entity having cardinality M
  - For total participation, use NOT NULL constraint

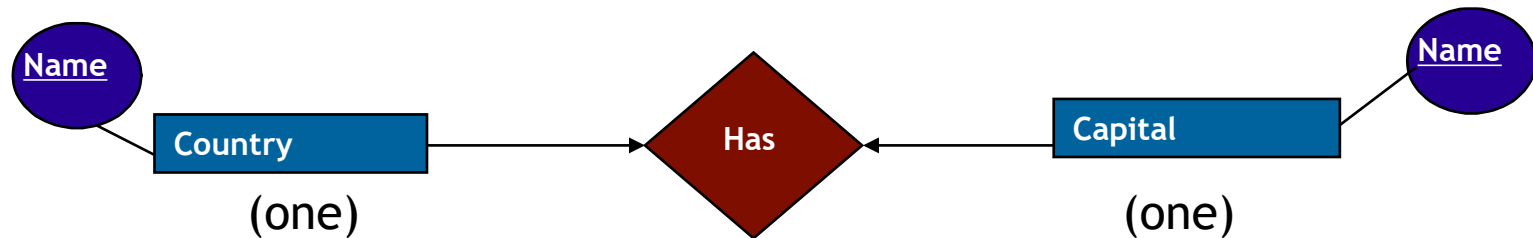
# Relationship Sets with Key Constraints (one to one case)

---



- Let's assume we went with Country(coName, caName) and all attributes have type Char(20) and we're not creating a separate relation for Capital. Write the SQL DDL that you would need for this relation.

# Relationship Sets with Key Constraints (one to one case)



- Let's assume we went with Country(coName, caName) and all attributes have type Char(20) and we're not creating a separate relation for Capital. Write the SQL DDL that you would need for this relation.

```
Create table country(
 coname char(20) Primary key,
 Caname char(20) unique NOT NULL
);
```

# Let see ... (SQL Server)

```
Create table country(
 coname char(20) Primary key,
 Caname char(20) unique
);
```

```
insert into country values ('Canada', NULL);
insert into country values ('India', 'New Delhi');
insert into country values ('Spain', NULL);
```

Violation of UNIQUE KEY constraint 'UQ\_\_country\_\_9EF576D51C2A0B95'.  
Cannot insert duplicate key in object 'dbo.country'. The duplicate  
key value is (<NULL>).  
The statement has been terminated.

|   | coname | Caname    |
|---|--------|-----------|
| 1 | Canada | NULL      |
| 2 | India  | New Delhi |

```
Create table country(
 coname char(20) Primary key,
 Caname char(20) unique not null,
);
```

```
insert into country values ('Spain', 'MADRID');
insert into country values ('India', 'New Delhi');
insert into country values ('Canada', 'OTTAWA');
insert into country values ('Spain', NULL);
insert into country values ('France', 'MADRID');
```

|   | coname | Caname    |
|---|--------|-----------|
| 1 | Spain  | MADRID    |
| 2 | India  | New Delhi |
| 3 | Canada | OTTAWA    |

# Let see... (Oracle)

```
create table country1(
 coname char(20) Primary key,
 Caname char(20) unique
);
```

```
insert into country1 values ('Canada', NULL);
insert into country1 values ('India', 'New Delhi');
insert into country1 values ('Spain', NULL);
```

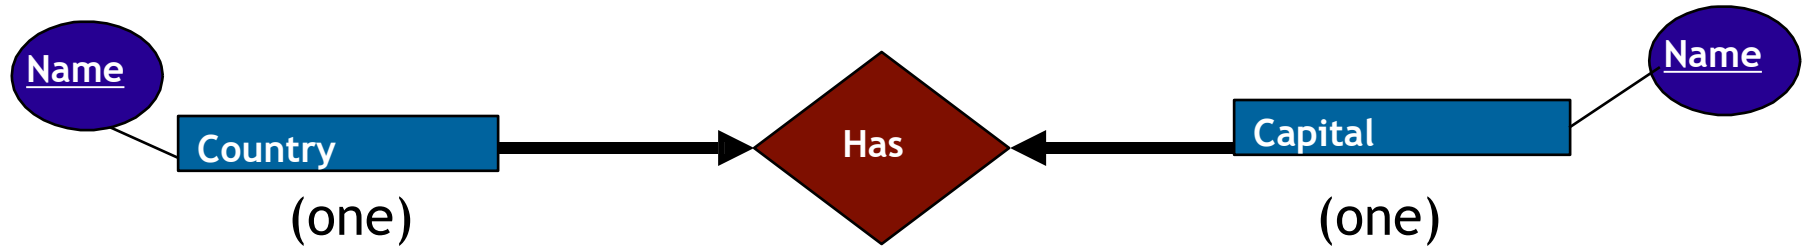
| CONAME | CANAME    |
|--------|-----------|
| Canada | -         |
| India  | New Delhi |
| Spain  | -         |

```
Create table country2(
 coname char(20) Primary key,
 Caname char(20) unique NOT NULL
);
```

```
insert into country2 values ('India', 'New Delhi');
insert into country2 values ('Canada', 'OTTAWA');
insert into country2 values ('Spain', 'MADRID');
insert into country2 values ('Spain', NULL);
insert into country2 values ('France', 'MADRID');
```

| CONAME | CANAME    |
|--------|-----------|
| India  | New Delhi |
| Canada | OTTAWA    |
| Spain  | MADRID    |

# Relationship Sets with Key Constraint + total participation

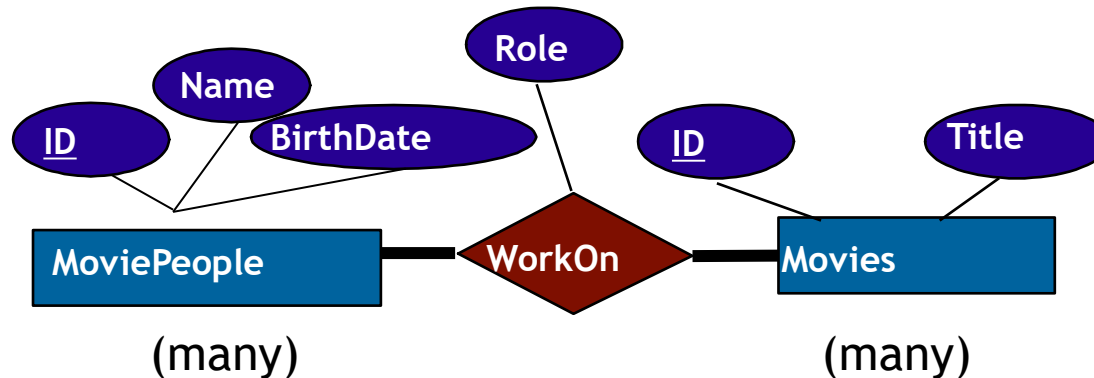


Let's assume we went with Country(coName, caName). Do we need a separate relation for Capital?

- A. Yes
- ☒ B. No
- C. It depends

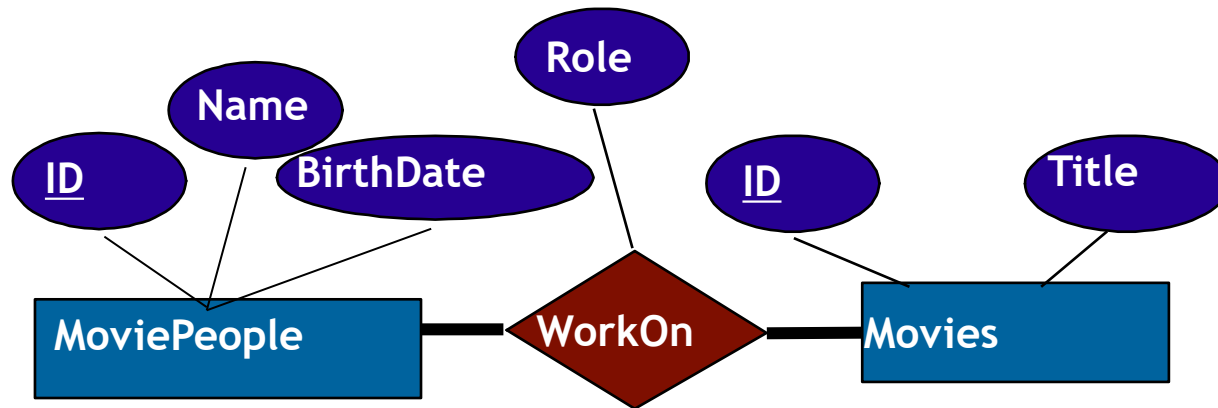
# Participation Constraints in SQL (cont')

- How can we express that “every movie person works on a movie and every movie has some movie person in it”?



- Neither foreign-key nor not-null constraints in **WorkOn** can do that.
- We need assertions (later)

# Let's see why we can't model this participation constraint using null constraint



MoviePeople

| <u>ID</u> | Name           | BirthDate  |
|-----------|----------------|------------|
| MP001     | Chris Columbus | 10/09/1958 |
| MP002     | Tom Hardy      | 15/09/1977 |
| MP003     | Kate Winslet   | 05/10/1975 |
| MP004     | Christian Bale | 30/01/1974 |

Movie

| <u>ID</u> | title        |
|-----------|--------------|
| 1         | Harry Potter |
| 2         | The Drop     |
| 3         | Legend       |
| 4         | Titanic      |

WorksOn

NOT NULL  
by default

| <u>PID</u> | <u>MID</u> | Role           |
|------------|------------|----------------|
| MP002      | 2          | Bob Saginowski |
| MP003      | 4          | Rose DeWitt    |

- We don't have all movies and all MoviePeople in the Works on. But its okay.
- Instance is a **legal** one!



# Glimpse of assertion

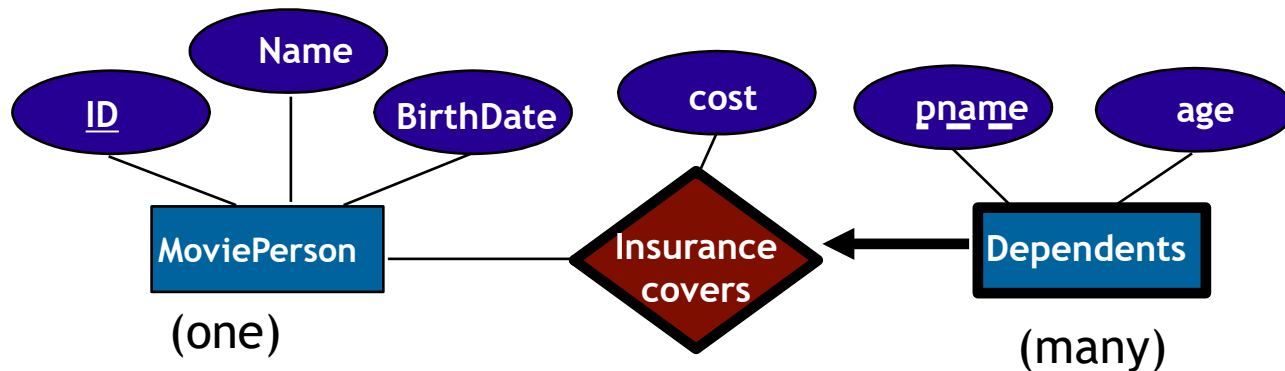
---

```
CREATE ASSERTION totalEmployment
CHECK
(NOT EXISTS ((SELECT PID FROM MoviePeople)
 EXCEPT
 (Select PID FROM WorkIn)));
```

# Translating Weak Entity Sets

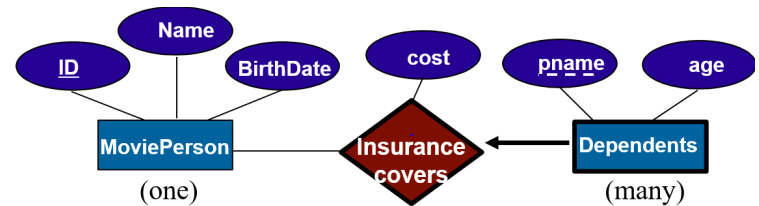
A **weak entity** is identified by considering the primary key of the *owner* (strong) entity.

- Owner entity set and weak entity set participate in a one-to-many identifying relationship set.
- Weak entity set has total participation.



- What is the best way to translate it?

# Translating Weak Entity Sets(cont')



- Weak entity set and its identifying relationship set are translated into a single table.
  - Primary key would consist of the owner's primary key and weak entity's partial key
  - When the owner entity is deleted, all owned weak entities must also be deleted.

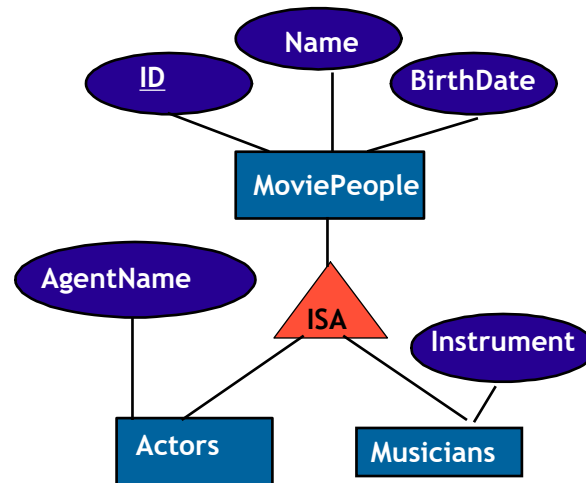
```
CREATE TABLE Dep_Insurance (
 pname CHAR(20),
 age INTEGER,
 cost REAL,
 ID CHAR(11)
 PRIMARY KEY (ID, pname),
 FOREIGN KEY (ID) REFERENCES MoviePeople
 ON DELETE CASCADE)
```

Same can be done in 1: M as well. If relationship has descriptive attribute, we can merge it. Just remember to have "CASCADING constraint" .

# Translating ISA Hierarchies to Relations

---

What is the best way to translate this into tables?



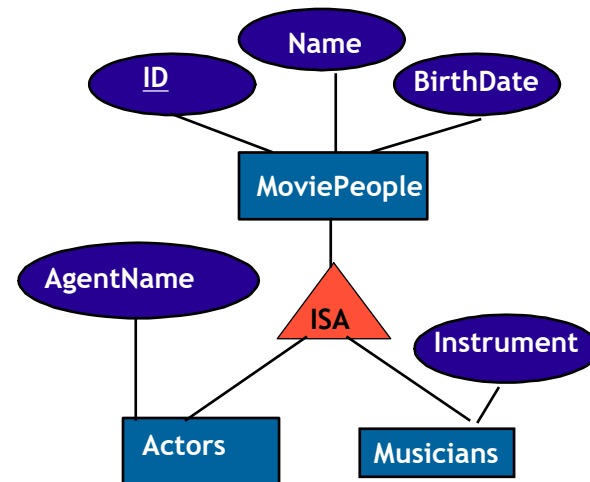
# Totally unsatisfactory attempt: Safest but with lots of duplication

One table per entity. Each has *all* attributes:

MoviePeople(ID, Name, BirthDate, AgentName, Instrument)

Actors(ID, Name, BirthDate, AgentName, Instrument)

Musicians(ID, Name, BirthDate, AgentName, Instrument)



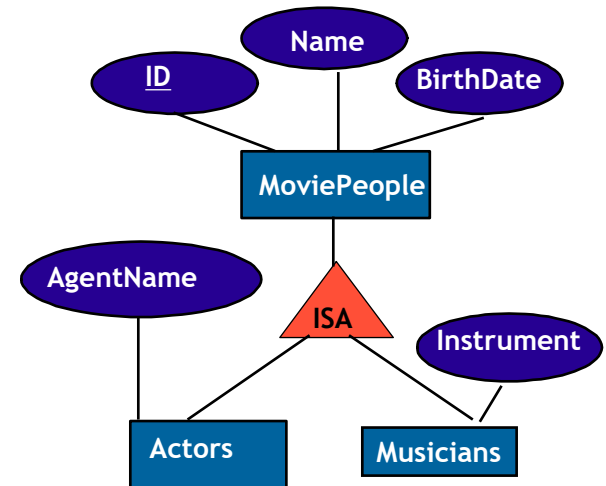
# Method 1: have only one table with *all* attributes

MoviePeople(ID, Name, BirthDate, AgentName, Instrument)

~~Actors(ID, Name, BirthDate, AgentName, Instrument)~~

~~Musicians(ID, Name, BirthDate, AgentName, Instrument)~~

- 1) What if I'm interested in just actors?
- 2) How to identify actors vs. musicians?
- 3) What if there is a relationship that only actors can participate in?
- 4) What if I wanted to add a new subclass



- Lots of space needed for nulls
- Excellent method if subclasses do not have any new attributes and relationships

## Method 2: 3 tables, remove excess attributes

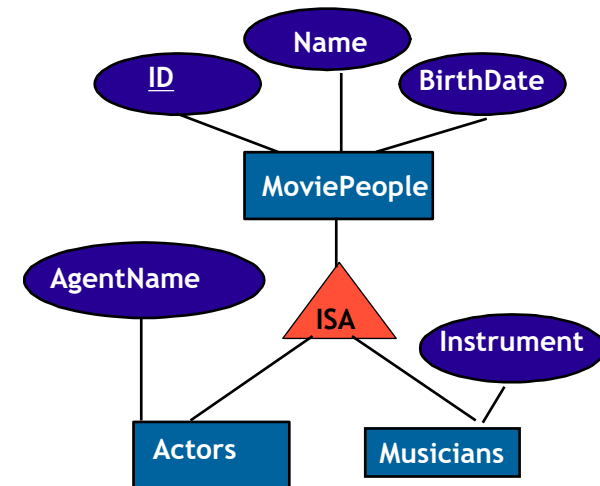
MoviePeople(ID, Name, BirthDate, ~~AgentName~~, ~~Instrument~~)

Actors(ID, ~~Name~~, ~~BirthDate~~, AgentName, ~~Instrument~~)

Musicians(ID, ~~Name~~, ~~BirthDate~~, ~~AgentName~~, Instrument)

- superclass table contains all superclass attributes
- subclass table contains primary key of superclass (as foreign key) and the subclass attributes

- Works well for concentrating on superclass.
- Have to combine two tables to get all attributes for a subclass



## Method 3: 2 tables, none for superclass

~~MoviePeople~~(ID, Name, BirthDate, AgentName, Instrument)

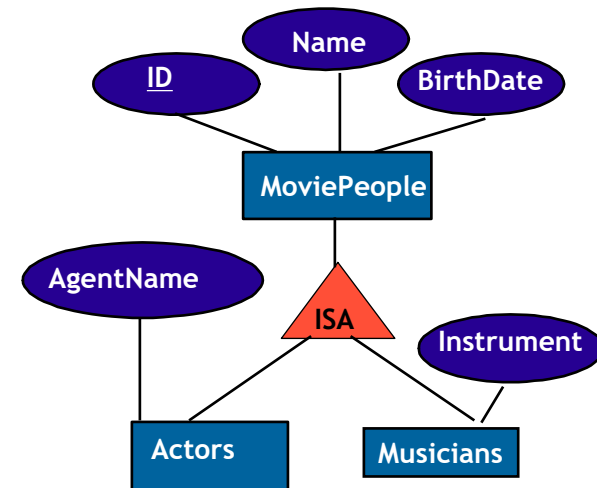
Actors(ID, Name, BirthDate, AgentName, ~~Instrument~~)

Musicians(ID, Name, BirthDate, ~~AgentName~~, Instrument)

- No table for superclass
- One table per subclass
- subclass tables have:
  - *all* superclass attributes
  - subclass attributes

How should we store directors?  
How do we store actors that are Musicians?

- Works poorly with relationships to superclass
- If ISA-relation is partial (not covering), it cannot be applied
- If ISA-relation is overlapping, it duplicates info





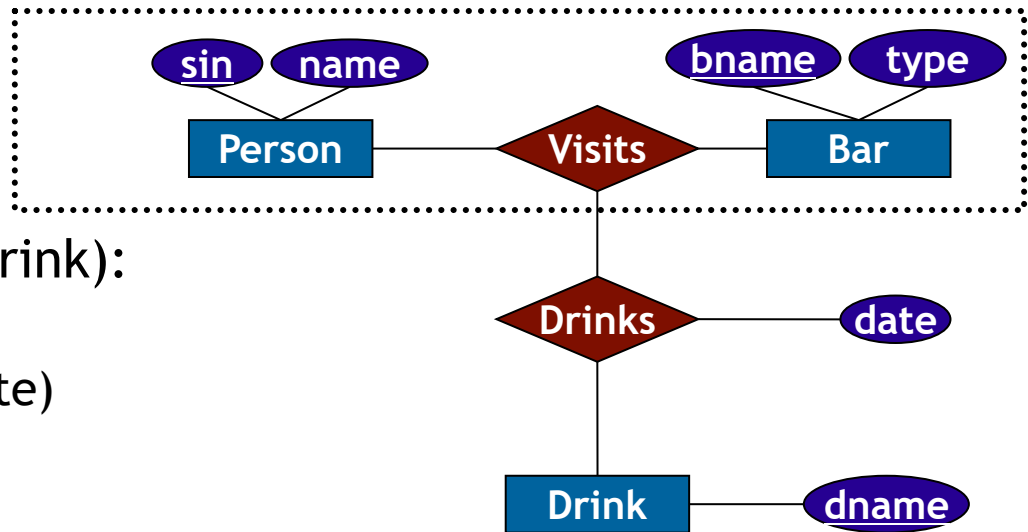
# Summary - ISA

---

- If subclasses have no new attributes or relationships, or if they all have the same attributes and relationships
  - One table with all attributes
- By default, ISA has **non-covering** and **disjoint** constraints
  - If ISA is covering and disjoint, subclasses have different attributes or different relationships and there is no need to keep the superclass
    - No table for the superclass
    - One table for each subclass with all attributes
- Otherwise (if ISA is non-covering or overlapping and subclasses have new and different attributes....)
  - A table for the superclass and one for each subclass

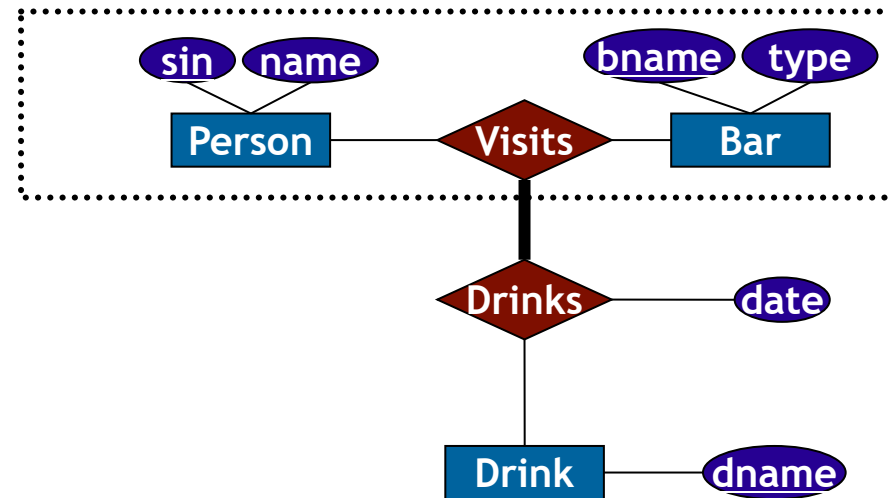
# Translating Aggregation

- Use standard mapping of relationship sets
- Tables for our example (other than Person, Bar, and Drink):
  - Visits(sin, bname)
  - Drinks(sin, bname, dname, date)



# Aggregation - Special Case

- There is a case where no table is required for an aggregate entity
  - Even where there are no cardinality constraints
    - If there is total participation between the aggregate entity and its relationship, and
    - If the aggregate entity does not have any descriptive attributes



If Visits is total on Drinks and Visits has no descriptive attributes we could keep only the Drinks

# Relational Model: Summary

---

- A tabular representation of data.
- Simple and intuitive, currently the most widely used.
- Integrity constraints can be specified, based on application semantics. DBMS checks for violations.
  - Important ICs: primary and foreign keys
  - Additional constraints can be defined with assertions (but are expensive to check)
- Powerful and natural query languages exist.
- Rules to translate ER to relational model